

OpenMP Offload Programming

Dr.-Ing. Michael Klemm
Chief Executive Officer
OpenMP Architecture Review Board



Agenda

- OpenMP Architecture Review Board
- ■Introduction to OpenMP Offload Features
- Case Study: NWChem TCE CCSD(T)
- Detachable Tasks



Introduction to OpenMP Offload Features



Running Example for this Presentation: saxpy

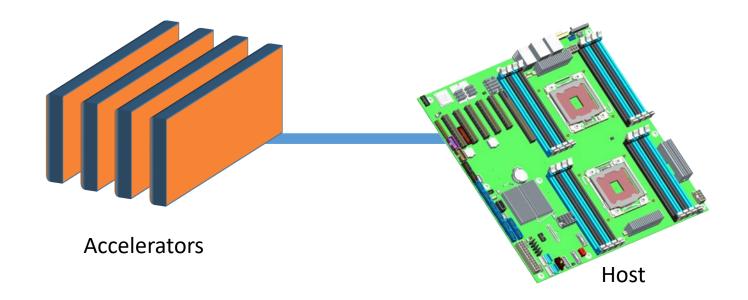
```
void saxpy() {
    float a, x[SZ], y[SZ];
    // left out initialization
    double t = 0.0;
                                                        Timing code (not needed, just to have
    double tb, te;
                                                        a bit more code to show (2)
    tb = omp_get_wtime();
#pragma omp parallel for firstprivate(a)
    for (int i = 0; i < SZ; i++) {
                                                        This is the code we want to execute on a
        y[i] = a * x[i] + y[i];
                                                        target device (i.e., GPU)
    te = omp_get_wtime();
                                                        Timing code (not needed, just to have
    t = te - tb;
                                                        a bit more code to show ①)
    printf("Time of kernel: %lf\n", t);
```

Don't do this at home!
Use a BLAS library for this!



Device Model

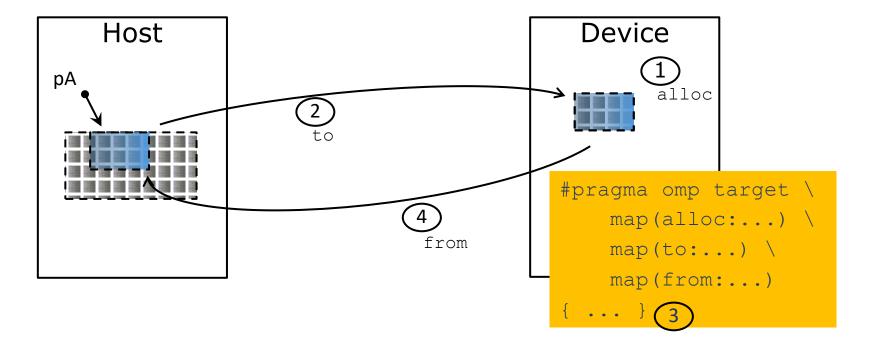
- As of version 4.0 the OpenMP API supports accelerators/coprocessors
- Device model:
 - One host for "traditional" multi-threading
 - Multiple accelerators/coprocessors of the same kind for offloading





Execution Model

- Offload region and data environment is lexically scoped
 - Data environment is destroyed at closing curly brace
 - Allocated buffers/data are automatically released





OpenMP for Devices - Constructs

- Transfer control and data from the host to the device
- Syntax (Fortran)

```
!$omp target [clause[[,] clause],...]
structured-block
!$omp end target
```

Clauses

```
device(scalar-integer-expression)
map([{alloc | to | from | tofrom}:] list)
if(scalar-expr)
```



```
The compiler identifies variables that are
                                                                    used in the target region.
void saxpy() {
    float a, x[SZ], y[SZ];
                                                                        All accessed arrays are copied from
    double t = 0.0;
                                                                             host to device and back
    double tb, te;
                                                             x[0:SZ]
    tb = omp get wtime();
                                                             y[0:SZ]
#pragma omp target "map(tofrom:y[0:SZ])"
    for (int i = 0; i < SZ; i++) {
         y[i] = a * x[i] + y[i];
                                                                             Presence check: only transfer
    te = omp_get_wtime();
                                                                              if not yet allocated on the
                                                            x[0:SZ]
    t = te - tb;
                                                            y[0:SZ]
                                                                                      device.
    printf("Time of kernel: %lf\n", t);
                                                                         Copying x back is not necessary: it
                                                                                was not changed.
```



The compiler identifies variables that are used in the target region.

```
subroutine saxpy(a, x, y, n)
    use iso_fortran_env
    integer :: n, i
                                                                     All accessed arrays are copied from
    real(kind=real32) :: a
                                                                          host to device and back
    real(kind=real32), dimension(n) :: x
                                                          x(1:n)
    real(kind=real32), dimension(n) :: y
                                                          y(1:n)
!$omp target "map(tofrom:y(1:n))"
                                                                         Presence check: only transfer
    do i=1,n
        y(i) = a * x(i) + y(i)
                                                                           if not yet allocated on the
    end do
                                                                                   device.
!$omp end target
                                                          x(1:n)
end subroutine
                                                          y(1:n)
```

Copying x back is not necessary: it was not changed.



```
void saxpy() {
    double a, x[SZ], y[SZ];
    double t = 0.0;
    double tb, te;
                                                      x[0:SZ]
    tb = omp_get_wtime();
                                                      y[0:SZ]
#pragma omp target map(to:x[0:SZ]) \
                   map(tofrom:y[0:SZ])
    for (int i = 0; i < SZ; i++) {
        y[i] = a * x[i] + y[i];
                                                     y[0:SZ]
   te = omp_get_wtime();
    t = te - tb;
    printf("Time of kernel: %lf\n", t);
```



```
The compiler cannot determine the size
                                                            of memory behind the pointer.
void saxpy(float a, float* x, float* y,
           int sz) {
    double t = 0.0;
    double tb, te;
                                                       x[0:sz]
    tb = omp_get_wtime();
                                                       y[0:sz]
#pragma omp target map(to:x[0:sz]) \
                    map(tofrom:y[0:sz])
    for (int i = 0; i < sz; i++) {
        v[0:sz]
    te = omp get wtime();
    t = te - tb;
    printf("Time of kernel: %lf\n", t);
                                                       Programmers have to help the compiler
                                                      with the size of the data transfer needed.
```



Creating Parallelism on the Target Device

- ■The target construct transfers the control flow to the target device
 - Transfer of control is sequential and synchronous
 - This is intentional!

- OpenMP separates offload and parallelism
 - Programmers need to explicitly create parallel regions on the target device
 - In theory, this can be combined with any OpenMP construct
 - In practice, there is only a useful subset of OpenMP features for a target device such as a GPU, e.g., no I/O, limited use of base language features.



Create a team of threads to execute the loop in parallel using SIMD instructions.



teams Construct

- ■Support multi-level parallel devices
- Syntax (C/C++):

```
#pragma omp teams [clause[[,] clause],...]
structured-block
```

■Syntax (Fortran):

```
!$omp teams [clause[[,] clause],...]
structured-block
```

■ Clauses

```
num_teams(integer-expression), thread_limit(integer-expression)
default(shared | firstprivate | private none)
private(list), firstprivate(list), shared(list), reduction(operator:list)
```



Multi-level Parallel saxpy

- Manual code transformation
 - Tile the loops into an outer loop and an inner loop
 - Assign the outer loop to "teams" (OpenCL: work groups)
 - Assign the inner loop to the "threads" (OpenCL: work items)



Multi-level Parallel saxpy

■ For convenience, OpenMP defines composite constructs to implement the required code transformations

```
subroutine saxpy(a, x, y, n)
   ! Declarations omitted
!$omp omp target teams distribute parallel do simd &
!$omp& num_teams(num_blocks) map(to:x) map(tofrom:y)
   do i=1,n
        y(i) = a * x(i) + y(i)
   end do
!$omp end target teams distribute parallel do simd
end subroutine
```



Optimize Data Transfers

- Reduce the amount of time spent transferring data
 - Use map clauses to enforce direction of data transfer.
 - Use target data, target enter data, target exit data constructs to keep data environment on the target device.

```
void zeros(float* a, int n) {
#pragma omp target teams distribute parallel for
   for (int i = 0; i < n; i++)
        a[i] = 0.0f;
}</pre>
```

```
void saxpy(float a, float* y, float* x, int n) {
#pragma omp target teams distribute parallel for
   for (int i = 0; i < n; i++)
      y[i] = a * x[i] + y[i];
}</pre>
```



target data Construct Syntax

- Create scoped data environment and transfer data from the host to the device and back
- Syntax (Fortran)

```
!$omp target data [clause[[,] clause],...]
structured-block
!$omp end target data
```

Clauses

```
device(scalar-integer-expression)
map([{alloc | to | from | tofrom | release | delete}:] list)
if(scalar-expr)
```



target update Construct Syntax

- Issue data transfers to or from existing data device environment
- Syntax (C/C++)

 #pragma omp target update [clause[[,] clause],...]
- Syntax (Fortran)
 !\$omp target update [clause[[,] clause],...]

Clauses

```
device(scalar-integer-expression)
to(list)
from(list)
if(scalar-expr)
```



Example: target data and target update

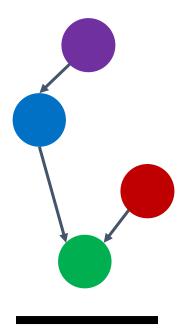
```
#pragma omp target data device(0) map(alloc:tmp[:N]) map(to:input[:N)) map(from:res)
#pragma omp target device(0)
#pragma omp parallel for
    for (i=0; i<N; i++)
      tmp[i] = some computation(input[i], i);
    update_input_array_on_the_host(input);
#pragma omp target update device(0) to(input[:N])
#pragma omp target device(0)
#pragma omp parallel for reduction(+:res)
    for (i=0; i<N; i++)
      res += final_computation(input[i], tmp[i], i)
```



Asynchronous Offloads

- OpenMP target constructs are synchronous by default
 - The encountering host thread awaits the end of the target region before continuing
 - The nowait clause makes the target constructs asynchronous (in OpenMP speak: they become an OpenMP task)

```
depend(out:a)
#pragma omp task
    init data(a);
                                                               depend(in:a) depend(out:x)
#pragma omp target map(to:a[:N]) map(from:x[:N])
                                                  nowait
    compute 1(a, x, N);
#pragma omp target map(to:b[:N]) map(from:z[:N])
                                                               depend(out:y)
                                                  nowait
    compute 3(b, z, N);
#pragma omp target map(to:y[:N]) map(to:z[:N])
                                                               depend(in:x) depend(in:y)
                                                   nowait
    compute 4(z, x, y, N);
#pragma omp taskwait
```





Case Study: NWChem TCE CCSD(T)

TCE: Tensor Contraction Engine

CCSD(T): Coupled-Cluster with Single, Double,

and perturbative Triple replacements



NWChem

- Computational chemistry software package
 - Quantum chemistry
 - Molecular dynamics
- Designed for large-scale supercomputers
- Developed at the EMSL at PNNL
 - EMSL: Environmental Molecular Sciences Laboratory
 - PNNL: Pacific Northwest National Lab
- URL: http://www.nwchem-sw.org



Finding Offload Candidates

- Requirements for offload candidates
 - Compute-intensive code regions (kernels)
 - Highly parallel
 - Compute scaling stronger than data transfer, e.g., compute O(n³) vs. data size O(n²)



Example Kernel (1 of 27 in total)

```
subroutine sd t d1 1(h3d,h2d,h1d,p6d,p5d,p4d,
                    h7d, triplesx, t2sub, v2sub)
     Declarations omitted.
     double precision triplesx(h3d*h2d,h1d,p6d,p5d,p4d)
     double precision t2sub(h7d,p4d,p5d,h1d)
     double precision v2sub(h3d*h2d,p6d,h7d)
!$omp target ,,presence?(triplesx,t2sub,v2sub)"
!$omp teams distribute parallel do private(p4,p5,p6,h2,h3,h1,h7)
     do p4=1,p4d
     do p5=1,p5d
                           1.5GB data transferred
     do p6=1,p6d
     do h1=1,h1d
                                (host to device)
     do h7=1,h7d
     do h2h3=1,h3d*h2d
      triplesx(h2h3,h1,p6,p5,p4)=triplesx(h2h3,h1,p6,p5,p4)
    1 - t2sub(h7,p4,p5,h1)*v2sub(h2h3,p6,h7)
     end do
     end do
                       1.5GB data transferred
     end do
                           (device to host)
     end do
     end do
     end do
!$omp end teams dis
                   parallel do
!$omp end target
     end subroutine
```

- All kernels have the same structure
- 7 perfectly nested loops
- Some kernels contain inner product loop (then, 6 perfectly nested loops)
- Trip count per loop is equal to "tile size" (20-30 in production)
- Naïve data allocation (tile size 24)
 - Per-array transfer for each target construct
 - triplesx: 1458 MB
 - t2sub, v2sub: 2.5 MB each



Invoking the Kernels / Data Management

Simplified pseudo-code

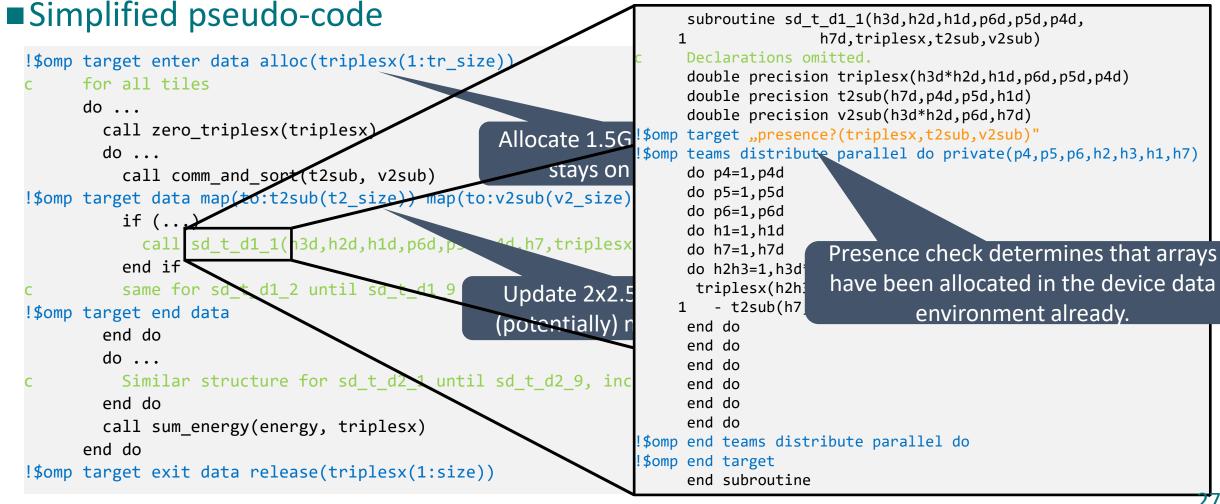
```
!$omp target enter data alloc(triplesx(1:tr_size))
     for all tiles
     do ...
       call zero triplesx(triplesx)
                                                Allocate 1.5GB data once,
       do ...
                                                      stays on device.
          call comm and sort(t2sub, v2sub)
!$omp target data map(to:t2sub(t2_size)) map(to:v2sub(v2_size))
          if (...)
           call sd_t_d1_1(h3d,h2d,h1d,p6d,p.
                                                1d_h7,triplesx,t2sub,v2sub)
          end if
          same for sd t d1 2 until sd t d1 9
                                                 Update 2x2.5MB of data for
!$omp target end data
                                                 (potentially) multiple kernels.
       end do
        do ...
          Similar structure for sd t d2 1 until sd t d2 9, incl. target data
       end do
        call sum_energy(energy, triplesx)
      end do
!$omp target exit data release(triplesx(1:size))
```

■ Reduced data transfers:

- triplesx:
 - allocated once
 - always kept on the target
- t2sub, v2sub:
 - allocated after comm.
 - kept for (multiple) kernel invocations



Invoking the Kernels / Data Management





Advanced Task Synchronization



Asynchronous API Interaction

- Some APIs are based on asynchronous operations
 - MPI asynchronous send and receive
 - Asynchronous I/O
 - HIP, CUDA and OpenCL stream-based offloading
 - In general: any other API/model that executes asynchronously with OpenMP (tasks)
- Example: CUDA memory transfers

```
do_something();
cudaMemcpyAsync(dst, src, nbytes, cudaMemcpyDeviceToHost, stream);
do_something_else();
cudaStreamSynchronize(stream);
do_other_important_stuff(dst);
```

- Programmers need a mechanism to marry asynchronous APIs with the parallel task model of OpenMP
 - How to synchronize completions events with task execution?



Try 1: Use just OpenMP Tasks

```
void cuda_example() {
#pragma omp task // task A
        do something();
        cudaMemcpyAsync(dst, src, nbytes, cudaMemcpyDeviceToHost, stream);
                                     Race condition between the tasks A & C,
    #pragma omp task // task B
                                      task C may start execution before
        do_something_else();
                                      task A enqueues memory transfer.
    #pragma omp task // task C
        cudaStreamSynchronize(stream);
        do other important stuff(dst);
```

■This solution does not work!



Try 2: Use just OpenMP Tasks Dependences

```
void cuda_example() {
#pragma omp task depend(out:stream) // task A
        do_something();
        cudaMemcpyAsync(dst, src, nbytes, cudaMemcpyDeviceToHost, stream);
                                                      Synchronize execution of tasks through dependence.
                                         // task B
    #pragma omp task
                                                      May work, but task C will be blocked waiting for
        do something else();
                                                      the data transfer to finish
    #pragma omp task depend(in:stream) // task C
        cudaStreamSynchronize(stream);
        do_other_important_stuff(dst);
```

- This solution may work, but
 - takes a thread away from execution while the system is handling the data transfer.
 - may be problematic if called interface is not thread-safe



OpenMP Detachable Tasks

- OpenMP 5.0 introduces the concept of a detachable task
 - Task can detach from executing thread without being "completed"
 - Regular task synchronization mechanisms can be applied to await completion of a detached task
 - Runtime API to complete a task

- Detached task events: omp event t datatype
- Detached task clause: detach(event)
- ■Runtime API: void omp_fulfill_event(omp_event_t *event)



Detaching Tasks

```
omp_event_t *event;
void detach_example() {
    #pragma omp task detach(event)
    {
        important_code();
    }
    #pragma omp taskwait ② ④
}

Some other thread/task:
    omp_fulfill_event(event);
```

- 1. Task detaches
- 2. taskwait construct cannot complete

- 3. Signal event for completion
- 4. Task completes and taskwait can continue



Putting It All Together

```
void CUDART CB callback(cudaStream t stream, cudaError t status, void *cb dat) {
 (3) omp_fulfill_event((omp_event_t *) cb_data);
void cuda example() {
    omp event t *cuda event;
#pragma omp task detach(cuda event) // task A
        do something();
        cudaMemcpyAsync(dst, src, nbytes, cudaMemcpyDeviceTbHost, stream);
        cudaStreamAddCallback(stream, callback, cuda event, 0);
#pragma omp task
                                     // task B
        do something else();
                                                         Task A detaches
#pragma omp taskwait(2)(4)
                                                      taskwait does not continue
#pragma omp task
                                     // task C
                                                      3. When memory transfer completes, callback is
                                                         invoked to signal the event for task completion
        do other important stuff(dst);
                                                      4. taskwait continues, task C executes
```



Removing the taskwait Construct

```
void CUDART_CB callback(cudaStream_t stream, cudaError_t status, void *cb_dat) {
 Omp_fulfill_event((omp_event_t *) cb_data);
void cuda_example() {
    omp event t *cuda event;
#pragma omp task depend(out:dst) detach(cuda event) // task A
        do something();
        cudaMemcpyAsync(dst, src, nbytes, cudaMemcpyDeviceTbHost, stream);
        cudaStreamAddCallback(stream, callback, cuda_event, 0);
#pragma omp task
                                     // task B
        do_something_else();
                                                          Task A detaches and task C will not execute because
                                                          of its unfulfilled dependency on A
#pragma omp task depend(in:dst)
                                                      2. When memory transfer completes, callback is
                                                          invoked to signal the event for task completion
        do other important stuff(dst);
```

Task A completes and C's dependency is fulfilled



Summary

- OpenMP API is ready to use Intel discrete GPUs for offloading compute
 - Mature offload model w/ support for asynchronous offload/transfer
 - Tightly integrates with OpenMP multi-threading on the host
- More, advanced features (not covered here)
 - Memory management API
 - Interoperability with native data management
 - Interoperability with native streaming interfaces
 - Unified shared memory support

OPENNIE 1997 Control of the control

Visit www.openmp.org for more information



Tools for OpenMP Programming



OpenMP Tools



- Correctness Tools
 - → ThreadSanitizer
 - →Intel Inspector XE (or whatever the current name is)

- Performance Analysis
 - → Performance Analysis basics
 - →Overview on available tools



Data Race



- Data Race: the typical OpenMP programming error, when:
 - >two or more threads access the same memory location, and
 - →at least one of these accesses is a write, and
 - →the accesses are not protected by locks or critical regions, and
 - →the accesses are not synchronized, e.g. by a barrier.
- Non-deterministic occurrence: e.g. the sequence of the execution of parallel loop iterations is non-deterministic
 - →In many cases *private* clauses, *barriers* or *critical regions* are missing
- Data races are hard to find using a traditional debugger



ThreadSanitizer: Overview



- Correctness checking for threaded applications
- Integrated in clang and gcc compiler
- Low runtime overhead: 2x 15x

Used to find data races in browsers like Chrome and Firefox



ThreadSanitizer: Usage

Module in Aachen.

https://pruners.github.io



module load clang



• Execute:

```
OMP_NUM_THREADS=4 ./myprog
```

Understand and correct the detected threading errors



C++

Fortran

ThreadSanitizer: Example

```
OpenMP
```

```
1 #include <stdio.h>
 3 int main(int argc, char **argv)
       int a = 0;
      #pragma omp parallel
 6
         if (a < 100) { <
            #pragma omp critical
            a++; ←
10
12 }
```

```
WARNING: ThreadSanitizer: data race
```

Read of size 4 at 0x7ffffffdcdc by thread T2:

```
#0 .omp_outlined. race.c:7
(race+0x0000004a6dce)
#1 __kmp_invoke_microtask <null>
(libomp_tsan.so)
```

Previous write of size 4 at 0x7fffffffdcdc by main thread:

```
#0 .omp_outlined. race.c:9
(race+0x0000004a6e2c)
#1 __kmp_invoke_microtask <null>
(libomp_tsan.so)
```



Intel Inspector XE



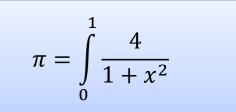
- Detection of
 - → Memory Errors
 - → Deadlocks
 - → Data Races
- Support for
 - →WIN32-Threads, Posix-Threads, Intel Threading Building Blocks and OpenMP
- Features
 - → Binary instrumentation gives full functionality
 - →Independent stand-alone GUI for Windows and Linux

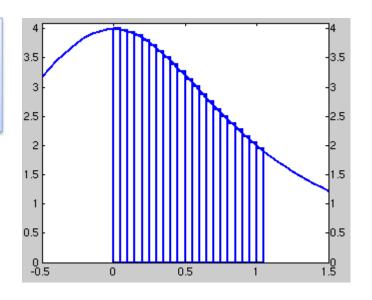


PI example / 1



```
double f(double x)
  return (4.0 / (1.0 + x*x));
double CalcPi (int n)
  const double fH = 1.0 / (double) n;
  double fSum = 0.0;
  double fX;
  int i;
#pragma omp parallel for private(fX,i) reduction(+:fSum)
  for (i = 0; i < n; i++)
    fX = fH * ((double)i + 0.5);
     fSum += \dot{f}(fX);
  return fH * fSum;
```







PI example / 2



```
double f(double x)
  return (4.0 / (1.0 + x*x));
double CalcPi (int n)
  const double fH = 1.0 / (double) n;
  double fSum = 0.0;
  double fX;
  int i;
#pragma omp parallel for private(fX,i) reduction(+:fSum)
  for (i = 0; i < n; i++)
    fX = fH * ((double)i + 0.5);
    fSum += f(fX);
  return fH * fSum;
```

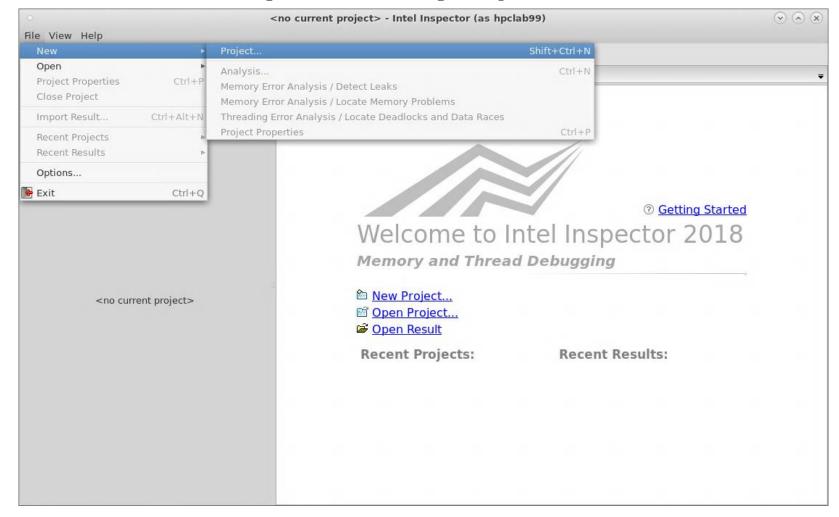
What if we would have forgotten this?



Inspector XE: create project / 1

Open**MP**

\$ module load Inspector ; inspxe-gui

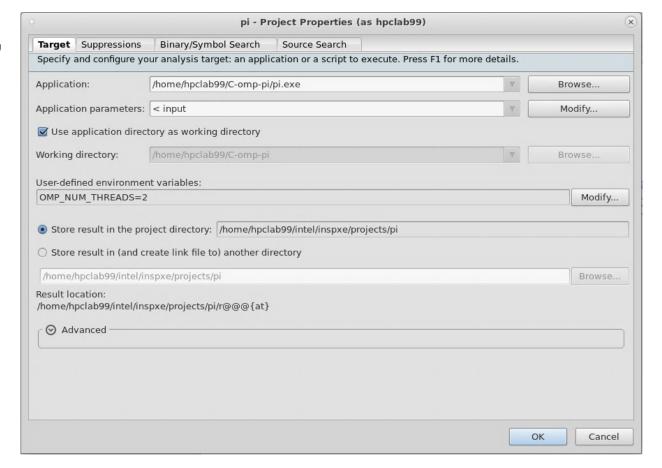




Inspector XE: create project / 2



- ensure that multiple threads are used
- choose a small dataset (really!),
 execution time can increase
 10X 1000X





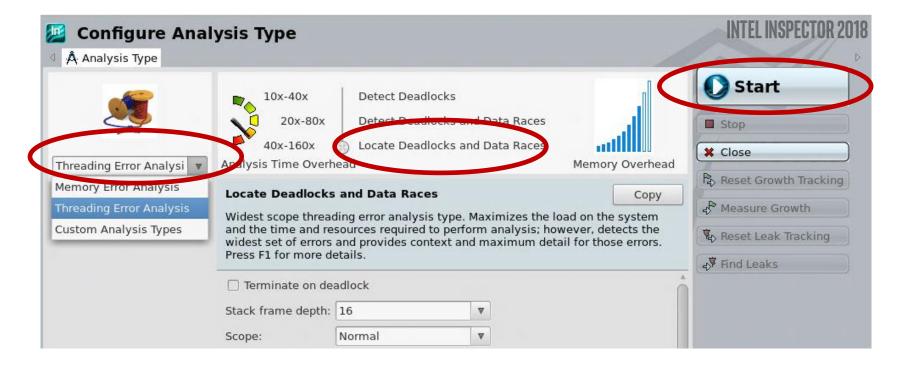
Inspector XE: configure analysis



Threading Error Analysis Modes

- Detect Deadlocks
- 2. Detect Deadlocks and Data Races
- 3. Locate Deadlocks and Data Races

more details, more overhead



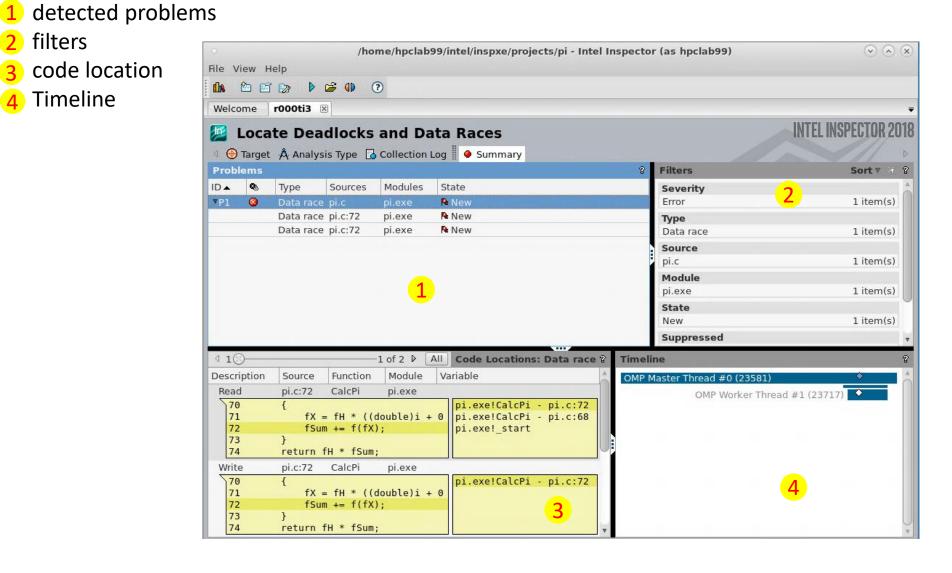


Inspector XE: results / 1

filters

code location

Timeline



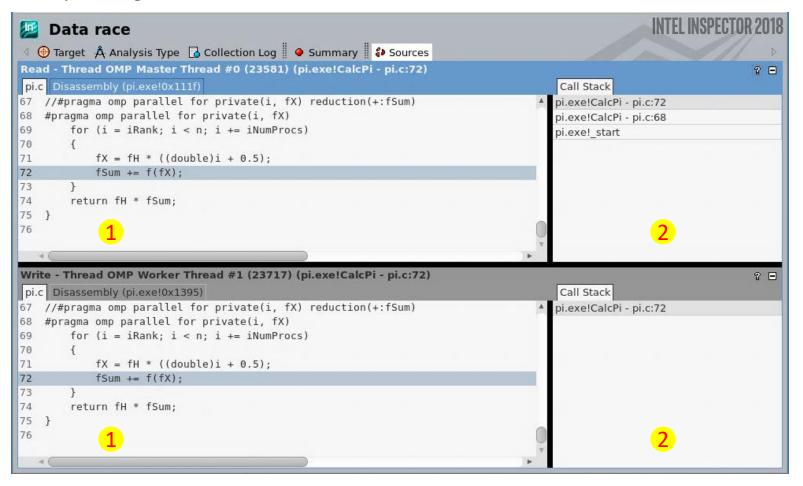




Inspector XE: results / 2

Open**MP**

- Source Code producing the issue double click opens an editor
- 2 Corresponding Call Stack





Inspector XE: results / 3

Open**MP**

Source Code producing the issue – double click opens an editor

Corresponding Call Stack The missing reduction Data race is detected. ⊕ Target A Analysis Type 🖟 Collection Log 🕴 🍑 Summary 🖁 🚱 Sources Read - Thread OMP Master Thread #0 (23581) (pi.exe!CalcPi - pi.c:72) 8 🖃 Disassembly (pi.exe!0x111f) Call Stack //#pragma omp parallel for private(i, fX) reduction(+:fSum) pi.exe!CalcPi - pi.c:72 #pragma omp parallel for private(i, fX) pi.exe!CalcPi - pi.c:68 for (i = iRank; i < n; i += iNumProcs) pi.exe! start fX = fH * ((double)i + 0.5);fSum += f(fX);72 73 74 return fH * fSum; 75 } Write - Thread OMP Worker Thread #1 (23717) (pi.exe!CalcPi - pi.c:72) Disassembly (pi.exe!0x1395) Call Stack 67 //#pragma omp parallel for private(i, fX) reduction(+:fSum) pi.exe!CalcPi - pi.c:72 #pragma omp parallel for private(i, fX) for (i = iRank; i < n; i += iNumProcs) fX = fH * ((double)i + 0.5);fSum += f(fX);72 73 return fH * fSum; 75 }

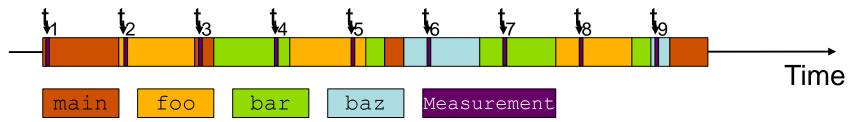


Sampling vs. Instrumentation



Sampling

- Running program is periodically interrupted to take measurement
- Statistical inference of program behavior
- Works with unmodified executables



Instrumentation

- Every event of interest is captured directly
- More detailed and exact information
- Typically: recompile for instrumentation



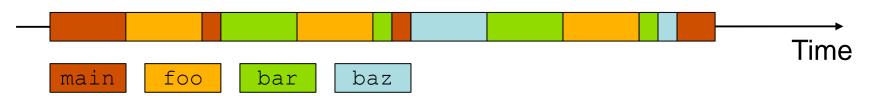


Tracing vs. Profiling



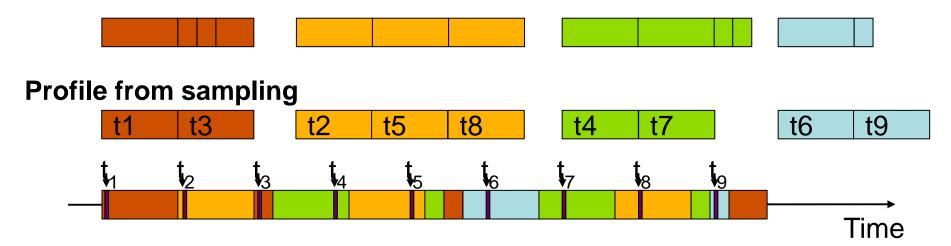
Trace

Chronologically ordered sequence of event records



Profile from instrumentation

Aggregated information





OMPT support for sampling

OMPT defines states like *barrier-wait*, *work-serial* or *work-parallel*

void foo() {}

int main()

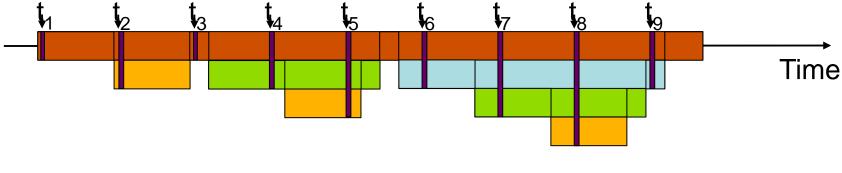
return 0;}

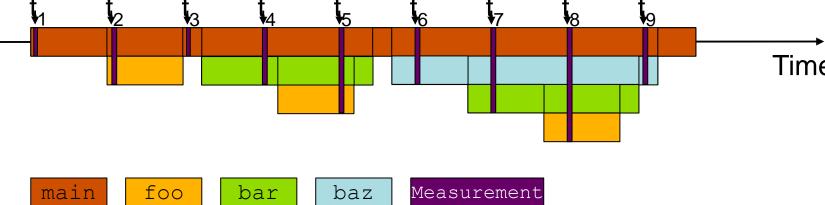
void bar() {foo();}

void baz() {bar();}

{foo();bar();baz();

- → Allows to collect OMPT state statistics in the profile
- → Profile break down for different OMPT states
- OMPT provides frame information
 - → Allows to identify OpenMP runtime frames.
 - → Runtime frames can be eliminated from call trees.







OMPT support for instrumentation



- OMPT provides event callbacks
 - → Parallel begin / end
 - → Implicit task begin / end
 - → Barrier / taskwait
 - → Task create / schedule
- Tool can instrument those callbacks
- OpenMP-only instrumentation might be sufficient for some use-cases

```
void foo() {}
void bar() {
    #pragma omp task
    foo();}
void baz() {
    #pragma omp task
       bar();}
int main() {
    #pragma omp parallel sections
    {foo();bar();baz();}
    return 0;}
```



VI-HPS Tools / 1



- Virtual institute high productivity supercomputing
- Tool development
- Training:
 - → VI-HPS/PRACE tuning workshop series
 - → SC/ISC tutorials
- Many performance tools available under vi-hps.org
 - → tools → VI-HPS Tools Guide
 - → Tools-Guide: flyer with a 2 page summary for each tool



VI-HPS Tools / 2

Open**MP**

Data collection

- Score-P: instrumentation based profiling / tracing
- Extrae: instrumentation based profiling / tracing

Data processing

Scalasca : trace-based analysis

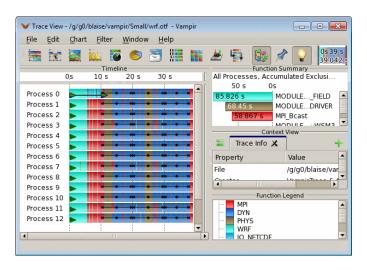
Data presentation

- ARM Map, ARM performance report
- CUBE : display for profile information
- Vampir : display for trace data (commercial/test)
- Paraver : display for extrae data
- Tau : visualization

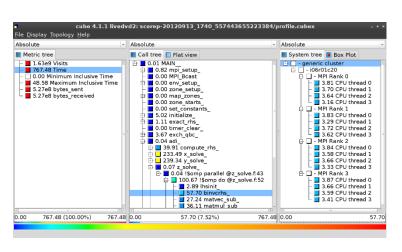


Performance tools GUI

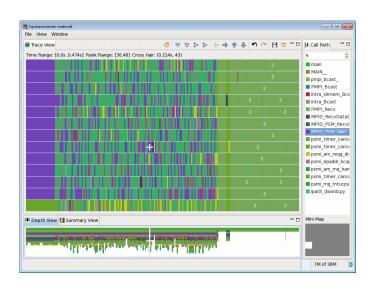












HPC Toolkit



Summary



Correctness:

- Data Races are very hard to find, since they do not show up every program run.
- Intel Inspector XE or ThreadSanitizer help a lot in finding these errors.
- Use really small datasets, since the runtime increases significantly.

Performance:

- Start with simple performance measurements like hotspots analyses and then focus on these hot spots.
- In OpenMP applications analyze the waiting time of threads. Is the waiting time balanced?
- Hardware counters might help for a better understanding of an application, but they might be hard to interpret.



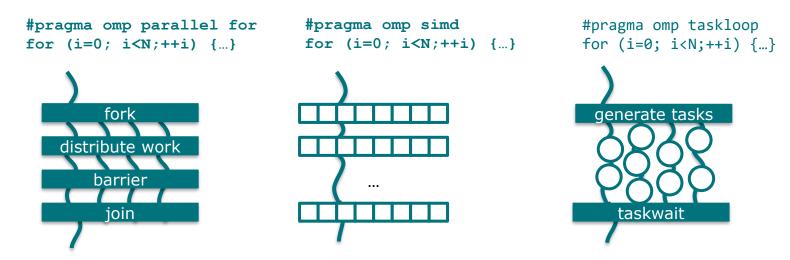
OpenMP Parallel Loops



100p Construct



Existing loop constructs are tightly bound to execution model:



The loop construct is meant to tell OpenMP about truly parallel semantics of a loop.



OpenMP Fully Parallel Loops



```
int main(int argc, const char* argv[]) {
    float *x = (float*) malloc(n * sizeof(float));
    float *y = (float*) malloc(n * sizeof(float));
    // Define scalars n, a, b & initialize x, y
#pragma omp parallel
#pragma omp loop
    for (int i = 0; i < n; ++i) {
     y[i] = a*x[i] + y[i];
```



loop Constructs, Syntax



Syntax (C/C++)

```
#pragma omp loop [clause[[,] clause],...]
for-loops
```

Syntax (Fortran)

```
!$omp loop [clause[[,] clause],...]
do-loops
[!$omp end loop]
```



loop Constructs, Clauses



- bind(binding)
 - → Binding region the loop construct should bind to
 - → One of: teams, parallel, thread
- order(concurrent)
 - → Tell the OpenMP compiler that the loop can be executed in any order.
 - → Default!
- \blacksquare collapse (n)
- private(list)
- lastprivate(list)
- reduction(reduction-id:list)



Extensions to Existing Constructs



Existing loop constructs have been extended to also have truly parallel semantics.

C/C++ Worksharing:

Fortran Worksharing:





DOACROSS Loops



DOACROSS Loops



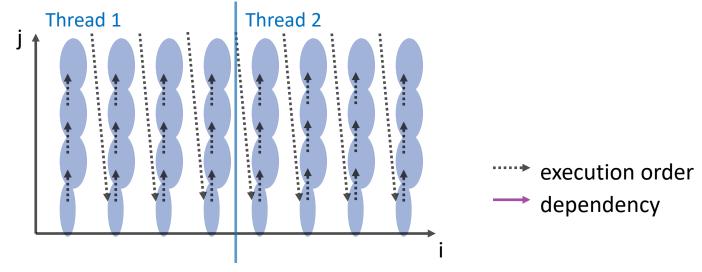
- "DOACROSS" loops are loops with special loop schedules
 - → Restricted form of loop-carried dependencies
 - → Require fine-grained synchronization protocol for parallelism
- Loop-carried dependency:
 - → Loop iterations depend on each other
 - → Source of dependency must scheduled before sink of the dependency
- DOACROSS loop:
 - → Data dependency is an invariant for the execution of the whole loop nest



Parallelizable Loops



A parallel loop cannot not have any loop-carried dependencies (simplified just a little bit!)

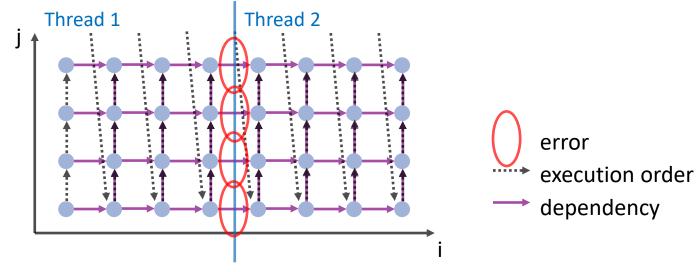




Non-parallelizable Loops



If there is a loop-carried dependency, a loop cannot be parallelized anymore ("easily" that is)

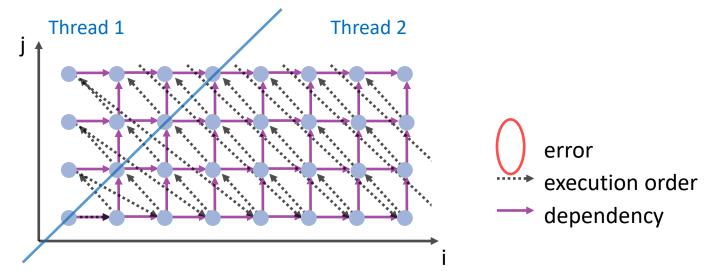




Wavefront-Parallel Loops



If the data dependency is invariant, then skewing the loop helps remove the data dependency





DOACROSS Loops with OpenMP



- OpenMP 4.5 extends the notion of the ordered construct to describe loop-carried dependencies
- Syntax (C/C++):

```
#pragma omp for ordered(d) [clause[[,] clause],...]
for-loops
and
#pragma omp ordered [clause[[,] clause],...]
where clause is one of the following:
    depend(source)
    depend(sink:vector)
```

Syntax (Fortran):

```
!$omp do ordered(d) [clause[[,] clause],...]
do-loops
!$omp ordered [clause[[,] clause],...]
```



Example



The ordered clause tells the compiler about loop-carried dependencies and their distances



Example: 3D Gauss-Seidel



```
#pragma omp for ordered(2) private(j,k)
for (i = 1; i < N-1; ++i) {
 for (j = 1; j < N-1; ++j)
#pragma omp ordered depend(sink: i-1,j-1) depend(sink: i-1,j) \
                    depend (sink: i-1, j+1) depend (sink: i, j-1)
   for (k = 1; k < N-1; ++k) {
      double tmp1 = (p[i-1][j-1][k-1] + p[i-1][j-1][k] + p[i-1][j-1][k+1]
                     + p[i-1][j][k-1] + p[i-1][j][k] + p[i-1][j][k+1]
                     + p[i-1][j+1][k-1] + p[i-1][j+1][k] + p[i-1][j+1][k+1]);
     double tmp2 = (p[i][j-1][k-1] + p[i][j-1][k] + p[i][j-1][k+1]
                     + p[i][j][k-1] + p[i][j][k] + p[i][j][k+1]
                     + p[i][j+1][k-1] + p[i][j+1][k] + p[i][j+1][k+1]);
      double tmp3 = (p[i+1][j-1][k-1] + p[i+1][j-1][k] + p[i+1][j-1][k+1]
                     + p[i+1][j][k-1] + p[i+1][j][k] + p[i+1][j][k+1]
                     + p[i+1][j+1][k-1] + p[i+1][j+1][k] + p[i+1][j+1][k+1]);
     p[i][j][k] = (tmp1 + tmp2 + tmp3) / 27.0;
#pragma omp ordered depend(source)
```





OpenMP API Version 5.1 State of the Union



Architecture Review Board

The mission of the OpenMP ARB (Architecture Review Board) is to standardize directive-based multi-language high-level parallelism that is performant, productive and portable.































































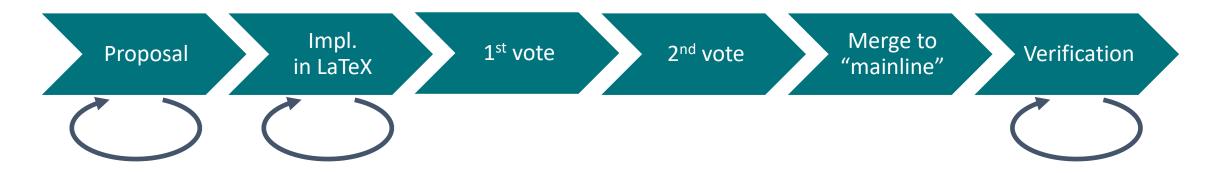






Development Process of the Specification

■ Modifications of the OpenMP specification follow a (strict) process:



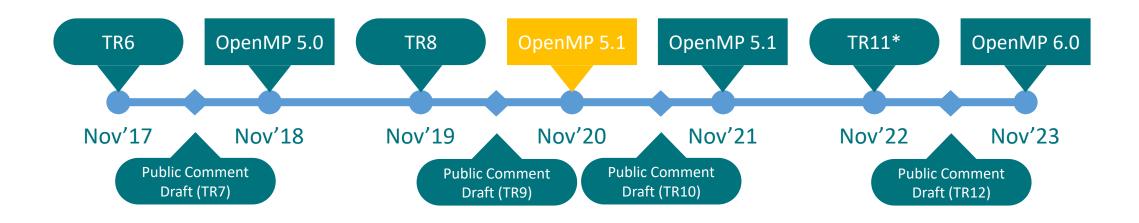
■ Release process for specifications:





OpenMP Roadmap

- OpenMP has a well-defined roadmap:
 - 5-year cadence for major releases
 - One minor release in between
 - (At least) one Technical Report (TR) with feature previews in every year



^{*} Numbers assigned to TRs may change if additional TRs are released.

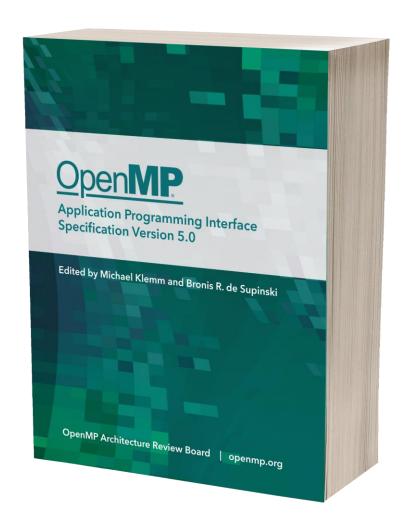


OpenMP API Version 6.0 Outlook – Plans

- Better support for descriptive and prescriptive control
- More support for memory affinity and complex memory hierarchies
- ■Support for pipelining, other computation/data associations
- Continued improvements to device support
 - Extensions of deep copy support (serialize/deserialize functions)
- Task-only, unshackled or free-agent threads
- Event-driven parallelism



Printed OpenMP API Specification

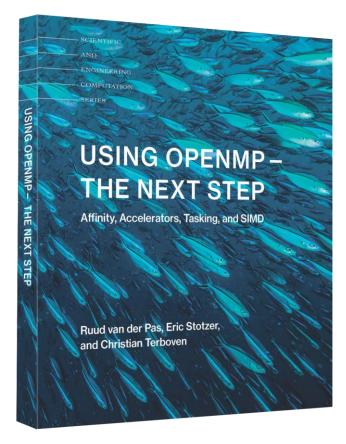


- Save your printer-ink and get the full specification as a paperback book!
 - Always have the spec in easy reach.
 - Includes the entire specification with the same pagination and line numbers as the PDF.
 - Available at a near-wholesale price.

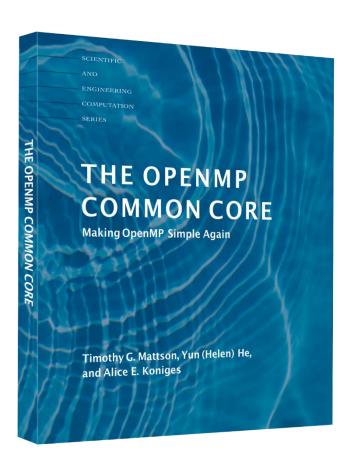
Get yours at Amazon at https://link.openmp.org/book51



Recent Books about OpenMP



Covers all of the OpenMP 4.5 features, 2017



Introduces the OpenMP Common Core, 2019



Help Us Shape the Future of OpenMP

- OpenMP continues to grow
 - 33 members currently

■ You can contribute to our annual releases

■ Attend IWOMP, become a cOMPunity member

- OpenMP membership types now include less expensive memberships
 - Please get in touch with me if you are interested

OPENNIE Enabling HPC since 1997

Visit www.openmp.org for more information